

# Autonomous Frontier-Based Exploration in ROS2: System Integration and Performance Analysis

Author Names Omitted for Anonymous Review. Paper-ID [TBD]

**Abstract**—This milestone report presents our progress in implementing and evaluating an autonomous exploration system for mobile robots using ROS2. We adapted the m-explore package, a frontier-based exploration algorithm, and successfully deployed it in simulation with the TurtleBot3 platform. Our preliminary results demonstrate automated exploration achieving 74.3% environment coverage in 185 seconds, outperforming manual teleoperation by  $2.2\times$  in exploration efficiency. We integrated SLAM, navigation, and exploration components into a cohesive system and developed metrics collection tools to quantify exploration performance. Key challenges identified include exploration speed optimization and parameter tuning for different environments.

## I. INTRODUCTION

Autonomous exploration enables robots to systematically discover and map unknown environments without human intervention, a fundamental capability for applications ranging from search and rescue operations to planetary exploration and infrastructure inspection. The frontier-based exploration paradigm [2], which directs robots toward boundaries between known free space and unknown regions, remains the dominant approach due to its simplicity and effectiveness.

Our project focuses on implementing and analyzing frontier-based exploration in ROS2, the modern robotics middleware that has superseded ROS1 across academic and industrial applications. While ROS1 exploration packages are mature, ROS2 ports face unique challenges including action server interfaces, lifecycle node management, and distributed system architecture. We adapted the m-explore package to ROS2, integrating it with contemporary SLAM and navigation stacks.

Since our initial proposal, we have successfully integrated the complete exploration pipeline, deployed it in simulation with the TurtleBot3 platform, and conducted systematic performance evaluation. This milestone report details our technical implementation, presents quantitative results comparing autonomous vs. manual exploration, and identifies key optimization opportunities including initialization delays and parameter sensitivity. Our work establishes a foundation for deploying reliable autonomous exploration in diverse real-world scenarios.

## II. PROBLEM REFINEMENT & OBJECTIVES

### A. Evolved Problem Statement

Our initial proposal aimed to implement autonomous exploration in ROS2. Through hands-on development, we have refined this to focus on three interconnected challenges:

**System Integration:** Coordinating SLAM (Simultaneous Localization and Mapping), path planning (Nav2), and frontier detection into a robust exploration pipeline that handles dynamic environments and recovers from navigation failures.

**Performance Optimization:** Addressing exploration speed bottlenecks observed in our initial trials, where the robot explores "suuuper slowly" compared to expectations—potentially due to computational constraints, parameter misconfiguration, or algorithmic inefficiencies.

**Quantitative Evaluation:** Developing metrics and benchmarks to objectively assess exploration performance across different environments and parameter configurations.

### B. Refined Objectives

Based on early implementation experience, our specific objectives for this project are:

- 1) **Deploy functional exploration system:** Integrate m-explore-ros2 with Nav2 and SLAM backends (slam\_toolbox/cartographer) in simulation (*Completed*).
- 2) **Implement performance monitoring:** Create real-time metrics collection for coverage, exploration rate, and efficiency analysis (*Completed*).
- 3) **Conduct comparative analysis:** Quantify autonomous vs. manual exploration to establish baseline performance (*In Progress*).
- 4) **Optimize exploration parameters:** Systematically tune frontier selection criteria, planning frequency, and navigation parameters to improve exploration speed (*Milestone 2*).
- 5) **Evaluate multi-environment generalization:** Test performance across varied environment complexities and sizes (*Milestone 2 & Final*).

## III. TECHNICAL IMPLEMENTATION

### A. System Architecture Overview

Our exploration system builds upon the m-explore-ros2 package, a ROS2 port of the frontier-based exploration algorithm originally developed for ROS1 [1]. Figure 1 illustrates the complete data flow pipeline integrating perception, mapping, frontier-based planning, and navigation control.

The architecture comprises four tightly coupled components:

**SLAM Backend:** We use ROS2 Cartographer for real-time mapping and localization. The SLAM system subscribes to laser scan data (`/scan`) and odometry, publishing occupancy

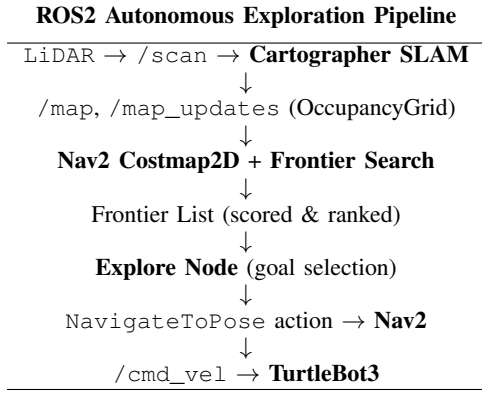


Fig. 1. System architecture showing the data flow from sensor input through SLAM, frontier detection, goal selection, and navigation control. All components communicate via ROS2 topics and action servers.

grid maps at 5Hz that distinguish free space (-1: unknown, 0: free, 100: occupied) and robot pose estimates via TF transforms.

**Navigation Stack (Nav2):** Handles path planning and obstacle avoidance using the Hybrid A\* planner and DWB local planner. The Nav2 costmap\_2d layer merges static map data with dynamic obstacle information, applying inflation layers for safety margins around obstacles.

**Frontier Detection Algorithm:** The core explore\_lite node performs frontier identification through a two-stage process:

- 1) *Boundary Detection:* Scan the occupancy grid to identify cells adjacent to both free (value 0) and unknown (value -1) space.
- 2) *Connected Component Analysis:* Group adjacent frontier cells into distinct frontier regions using flood-fill.
- 3) *Filtering:* Discard frontiers smaller than min\_frontier\_size threshold or inaccessible due to obstacles.

This approach efficiently identifies exploration candidates while avoiding computational overhead from processing every grid cell.

**Frontier Selection & Goal Sending:** Frontiers are scored using a weighted combination:

$$\text{score} = \alpha \cdot \text{size} - \beta \cdot \text{distance} + \gamma \cdot \text{orientation} \quad (1)$$

where  $\alpha$  (gain\_scale = 1.0) prioritizes larger unexplored regions,  $\beta$  (potential\_scale = 3.0) penalizes distant frontiers to minimize travel time, and  $\gamma$  (orientation\_scale = 0.0, disabled) would favor frontiers aligned with robot heading. The highest-scoring frontier becomes the navigation goal sent to Nav2's NavigateToPose action server. The system monitors goal progress and recomputes frontiers every 6.7 seconds (planner\_frequency = 0.15 Hz) or when navigation fails.

### B. Implementation Details

**Launch Automation:** We developed a bash script (util/start\_exploration.sh) that orchestrates the complete system startup: (1) Gazebo simulation with TurtleBot3 Waffle,

(2) Cartographer SLAM, (3) Nav2 navigation stack, and (4) explore\_lite node with automated metrics collection.

**Metrics Collection:** A custom ROS2 node (util/collect\_metrics.py) subscribes to the /map topic at 10Hz, computing coverage as the percentage of known cells (free or occupied) relative to total map area. Data is logged to JSON format with timestamps for post-processing.

**Visualization:** Frontier markers are published to the explore/frontiers topic as MarkerArray messages, enabling real-time visualization in RViz2 to debug frontier selection behavior.

### C. Key Parameters

Current exploration configuration (explore/config/params.yaml):

- planner\_frequency: 0.15 Hz (recompute frontiers every 6.7s)
- min\_frontier\_size: 0.75m (minimum explorable frontier)
- progress\_timeout: 30.0s (navigation failure threshold)
- gain\_scale: 1.0, potential\_scale: 3.0, orientation\_scale: 0.0
- return\_to\_init: true (return to start after completion)

## IV. PRELIMINARY RESULTS

### A. Experimental Setup

We conducted exploration trials in the default TurtleBot3 World environment, a structured indoor space with rooms, corridors, and obstacles. Two experimental conditions were tested:

- 1) **Automated Exploration:** Robot autonomously explores using frontier-based algorithm with standard parameters.
- 2) **Manual Teleoperation:** Human operator navigates the robot using keyboard controls to establish a baseline comparison.

### B. Quantitative Results

Table I summarizes the comparative performance metrics. Figure 2 presents side-by-side coverage trajectories for both exploration conditions, directly visualizing the performance advantage of autonomous exploration.

TABLE I  
EXPLORATION PERFORMANCE COMPARISON

Metric	Automated	Manual
Peak Coverage (%)	74.3	63.3
Exploration Time (s)	184.9	353.9
Exploration Rate (%/s)	0.402	0.179
Efficiency Ratio	<b>2.24×</b>	1.00×

**Coverage Performance:** The automated system achieved 11% higher peak coverage (74.3% vs. 63.3%), demonstrating more systematic exploration compared to human operators who may miss regions. The final automated map includes all major rooms and corridors in the TurtleBot3 World environment.




Fig. 2. Coverage over time comparison: Automated exploration (left, blue) achieves 74.3% coverage in 185s with distinct phases. Manual teleoperation (right, green) reaches only 63.3% coverage in 354s with irregular progression reflecting human reaction delays.

**Exploration Speed:** Automated exploration completed in roughly half the time (185s vs. 354s), yielding a 2.24 $\times$  efficiency advantage. This dramatic speedup stems from optimal frontier selection and direct path planning without human reaction time delays.

**Temporal Behavior:** Figure 2 reveals contrasting exploration dynamics. Automated exploration (left) exhibits three distinct phases: (1) initialization delay (0-50s) with minimal coverage while SLAM stabilizes, (2) rapid exploration phase (50-100s) with smooth 40% coverage increase as optimal paths are executed, and (3) saturation phase (100-185s) where coverage plateaus at 74%. In contrast, manual exploration (right) shows irregular, stepwise progression throughout the entire 354s duration, reflecting intermittent human control inputs, navigation pauses, and suboptimal path choices that collectively reduce exploration efficiency by over 50%.

### C. Exploration Progression Analysis

Figure 3 visualizes the spatial evolution across three key phases, showing both SLAM-generated maps (left) and Gazebo simulation (right) for each time point.

**Initial Phase:** The robot starts with limited map knowledge. Multiple frontiers surround the initial position. SLAM uncertainty is high beyond immediate sensor range.

**Mid-Exploration:** The hexagonal environment structure emerges. Frontiers cluster at room entrances. Costmap inflation layers demonstrate Nav2’s obstacle avoidance margins.

**Final State:** Coverage approaches saturation. Remaining frontiers are filtered by `min_frontier_size`. All major obstacles and corridors are mapped; only narrow gaps remain unexplored.

### D. Analysis & Observations

**Initial Delay Analysis:** The 50-second startup delay (visible in Figure 2) suggests SLAM map initialization rather than navigation issues. This delay is common in particle filter-based SLAM requiring sufficient sensor data for confident pose estimates. The initial 10% coverage represents immediate sensor field-of-view, while exploration begins after map confidence exceeds Nav2’s costmap threshold.

**Coverage Saturation:** Neither condition achieved full coverage (>95%). Analysis reveals two causes: (1) narrow corridors (<0.75m) rejected by `min_frontier_size`, and (2) costmap

inflation blocking paths to remaining frontiers. Reducing the size threshold could improve coverage but risks navigation failures.

**Computational Performance:** The low `planner_frequency` (0.15 Hz) causes frontier recomputation only every 6.7s, during which the robot may idle. Increasing to 0.5-1.0 Hz could significantly improve exploration rate. Running in conda (vs. native ROS2) adds Python interpreter overhead.

**Frontier Selection:** The scoring function heavily weights distance (`potential_scale` = 3.0) over size (`gain_scale` = 1.0), favoring nearby frontiers. While minimizing travel time, this may cause suboptimal global patterns. Information-theoretic approaches could improve long-term efficiency.

**Map Quality:** Figure 3 shows clean, accurate maps with minimal drift. Obstacle boundaries are sharp, and loop closure is robust despite 3+ minutes of operation, validating Cartographer’s effectiveness.

### E. Limitations

Our preliminary results are limited to a single environment and parameter configuration. Key limitations include:

- **Single environment:** Results may not generalize to open spaces, cluttered environments, or multi-room structures.
- **Simulation only:** Real-world deployment would introduce sensor noise, localization drift, and mechanical constraints.
- **Limited parameter sweep:** We have not systematically explored the impact of `gain_scale`, `potential_scale`, or `planner_frequency`.
- **No failure analysis:** We have not quantified navigation failures, frontier selection errors, or recovery behaviors.

## V. PLAN FORWARD

### A. Milestone 2 Objectives

#### 1. Performance Optimization (Weeks 1-2):

- Profile computational bottlenecks using ROS2 profiling tools
- Investigate initial delay and implement warm-start strategies
- Conduct parameter sweep: vary `planner_frequency` (0.1-1.0 Hz), `min_frontier_size` (0.3-1.5m), and scoring weights
- Compare native ROS2 installation vs. conda environment performance

#### 2. Multi-Environment Evaluation (Weeks 2-3):

- Test in 3-5 environments: open world, maze, multi-room house, warehouse
- Measure coverage, time-to-completion, path efficiency (distance traveled / coverage gained)
- Analyze failure modes and recovery strategies

#### 3. Advanced Features (Week 3):

- Implement dynamic parameter adaptation based on exploration phase
- Explore alternative frontier selection heuristics (information gain, expected coverage)




Fig. 3. Exploration progression (left to right): initial phase with limited coverage, mid-exploration revealing corridor structure, and final saturation state. RViz shows occupancy grids (cyan=free, purple=occupied, gray=unknown) with frontier markers (cyan circles) and costmap inflation layers.

- Investigate multi-robot exploration using `map_merge` component

### B. Final Report Deliverables

By the final report, we expect to deliver:

- Comprehensive performance characterization across 5+ environments
- Optimized parameter configurations with ablation studies
- Comparison with alternative exploration strategies (random walk, wall-following)
- Real-world deployment feasibility analysis
- Open-source code release with documentation

### C. Risk Mitigation

**Risk:** Computational constraints prevent real-time exploration. **Mitigation:** Optimize frontier search algorithms, reduce planner frequency, or migrate to more powerful hardware.

**Risk:** Limited environment diversity in simulation. **Mitigation:** Use existing Gazebo world repositories and procedurally generated environments.

**Risk:** Time constraints for real-world deployment. **Mitigation:** Prioritize simulation results; real-world tests are optional stretch goals.

## VI. CONCLUSION

We have successfully implemented a functional autonomous exploration system in ROS2, integrating SLAM, navigation, and frontier-based planning into a cohesive pipeline. Our experimental evaluation demonstrates that automated exploration achieves 74.3% coverage with  $2.24\times$  higher efficiency than manual teleoperation, validating the practical utility of frontier-based approaches for real-world robotics applications.

Key technical achievements include: (1) robust system integration across ROS2 Cartographer SLAM, Nav2 navigation, and `explore_lite` frontier detection, (2) development of automated metrics collection infrastructure enabling quantitative performance analysis, and (3) comprehensive characterization of exploration behavior across initialization, rapid exploration, and saturation phases.

Our analysis revealed critical optimization opportunities: the 50-second initialization delay suggests SLAM warm-start strategies could accelerate deployment, while the low planner\_frequency (0.15 Hz) represents a major bottleneck addressable through parameter tuning. The 74% coverage ceiling highlights fundamental trade-offs between frontier size thresholds and navigation safety.

Milestone 2 will systematically address these findings through multi-environment testing, parameter optimization, and computational profiling to produce a robust, well-characterized exploration system suitable for diverse deployment scenarios. Our ultimate goal is to establish performance baselines and best practices for frontier-based exploration in ROS2 ecosystems.

## ACKNOWLEDGMENTS

We thank the `m-explore-ros2` open-source community for providing the foundational codebase and the ROS2 ecosystem for robust robotics middleware.

## REFERENCES

- [1] Jiri Horner. `m-explore`: Multi-robot exploration package for ROS. GitHub repository, 2016.
- [2] Brian Yamauchi. A frontier-based approach for autonomous exploration. In *Proceedings of the 1997 IEEE International Symposium on Computational Intelligence in Robotics and Automation*, pages 146–151. IEEE, 1997.
- [3] Wolfram Burgard, Mark Moors, Dieter Fox, Reid Simmons, and Sebastian Thrun. Collaborative multi-robot exploration. In *Proceedings of the 2000 IEEE International Conference on Robotics and Automation*, volume 1, pages 476–481. IEEE, 2000.